

# OS HW 3 Dry part

Submitters:

Shaked Or

308026400

[shakedor@campus.technion.ac.il](mailto:shakedor@campus.technion.ac.il)

Aviad Rozenkof

316592922

[rozen501@gmail.com](mailto:rozen501@gmail.com)

Q1:

1. Let us label a linked list of size  $n$  as  $L$  and the elements of said list as  $\{a_i \mid 1 \leq i \leq n\}$ .

Note that each element has a mutex lock in its node.

Searching:

We lock  $a_1$  and then look at the data of  $a_1$ , if  $a_1$  does not contain the data desired we lock  $a_2$  unlock  $a_1$  (after which only  $a_2$  is locked by us) and then look at the data of  $a_2$ .

We continue to do so until the end of the list or until our element is found. If found we keep the lock only on the desired element and if endOfList is reached we unlock the last node before declaring failure. The idea is that we lock  $a_i$  and  $a_{i+1}$  when moving from the former to the latter so that the connection between them cannot be changed while we attempt to move between them.

Inserting: suppose we want to insert  $a_j$  between  $a_i$  and  $a_k = a_{i+1}$ . We lock  $a_j$  and  $a_k$  in that order to avoid deadlocking (unless the convention of the program is to iterate backwards). We then create and lock  $a_j$  (not necessary if no other thread knows  $a_j$ ) and link it to  $a_i$  and  $a_k$  (note that during the linking all 3 are locked by our thread). We then unlock all 3 nodes in any order.

Removing: suppose we want to remove  $a_j$  from in between  $a_i$  and  $a_k$  ( $i=j-1, k=j+1$ ).

We acquire the locks of  $a_i, a_j$  and  $a_k$  in that order to avoid deadlocking. Then we delink  $a_j$  from  $a_i$  and  $a_k$  and link  $a_i$  and  $a_k$  to each other. Then we can deallocate  $a_j$ . Afterwards, we unlock, the now linked,  $a_i$  and  $a_k$  in any order we choose.

Note that when inserting and removing element to and from the 2 ends of the list (meaning first and last) we do not have a previous/next element to lock and thus we only lock and link the subset of  $\{L_{neighbor}, target, R_{neighbor}\}$  that exist in the list.

2. Observation 1:

Let there be  $T_1, p_1, T_2$  and  $p_2$  as given in Observation 1 (denoted obs1). Since  $T_2$ 's candidate,  $p_2$ , is bigger than  $p_1$ , this means that  $T_2$  tried to select  $p_1$  as a candidate before but upon not finding  $p_1^2$  realized that another thread is handling  $p_1$ . This means that at that point in time,  $T_1$  was ahead of  $T_2$  in the list (since it was handling  $p_1$  before  $T_2$ ). Then  $T_2$  went back to  $p_1$  and tried getting another candidate and eventually reached  $p_2$ . Since we assume that  $T_1$  had not finished removing elements that are divisible by  $p_1$ , and we know that a thread cannot overtake another thread that is ahead of it in the list (unless the leading thread goes back to choose a new candidate) then  $T_2$  is necessarily behind  $T_1$  in the list in the scenario described by obs1. Thus the element locked by  $T_2$  is smaller than the element locked by  $T_1$ .

Observation 2:

Reminder, a thread  $T$  with a given candidate  $p$  can only remove elements that are divisible by  $p$ . Let  $p$  be a prime candidate. By obs3 (explained later)  $p$  is prime. Thus  $p^2$  is only divisible by 1,  $p$  and itself. 1 is not in the list and thus is never a candidate. If  $p^2$  is a candidate it is a prime by obs3 which contradicts the fact that it is the square of a prime. Thus the only possible thread that can remove it is  $T$  with given candidate  $p$ .

Observation 3:

Let there be  $p$  a prime candidate of thread  $T$ . If  $T$  managed to reach the state that it chose  $p$  as a candidate, that means that for all  $k < p$  either  $k$  was removed by a thread prior to this moment (possibly by  $T$  itself) or  $k$  was a candidate that  $T$  must have considered, meaning it either chose  $k$

as a candidate and removed all of the elements corresponding to  $k$  or it reached  $k^2$  and realized that another thread was handling it.

Let us assume that  $p$  is not prime. Thus it has a minimal prime component,  $e$ . Since  $e$  is prime and is smaller than  $p$ ,  $T$  must have considered  $e$  as a candidate. This leaves us with 3 cases:

- $T$  handled  $e$
- Another thread  $T'$  handled  $e$
- Another thread  $T'$  is currently handling  $e$

From the first 2 and the fact that  $p$  is divisible by  $e$  we get a contradiction since  $p$  was removed from the list and yet was selected as a candidate.

From the 3<sup>rd</sup> and obs1 we can deduce that the element locked by  $T'$ ,  $e'$ , is bigger than  $p$ . meaning that  $T'$  already passed  $p$  but since  $e$  divides  $p$  then  $p$  was removed. Contradiction.

Thus  $p$  is a prime candidate.

3. If all threads run the unmodified algorithm, then a thread  $T$  will reach candidate  $p$  only after reaching the end of the list trying to remove divisors of  $k$  for every  $k$ , a candidate smaller than  $p$ . Now, let  $T$  be the thread that selected 2 as a candidate first and  $T'$  any other thread.

We will show by induction that  $T'$  could not have removed any element.

Since all threads run the unmodified algo, then  $T'$  will only choose 3 as a candidate after it had reached the end of the list trying to find elements divisible by 2. Note that it won't find a single element divisible by 2 since it cannot overtake  $T$  and  $T$  is removing all of those elements. But if  $T'$  reached the end of the list it means that  $T$  already reached there and chose 3 as its candidate first. Meaning, like the case with 2, that  $T$  will remove all divisors of 3 that remain and  $T'$  will remove nothing. This goes on by induction for every candidate. Thus, the thread that reached 2 first will remove all elements and the other threads will do no real work.

Q2:

1. All nodes in the list contain a read write lock. All threads use hand over hand read/write locking. The threads will use read locks to search the list and acquire write locks when trying to remove elements from the list. When a thread sees a candidate is already being read it assumes another thread is handling that candidate and moves on to the next candidate. Under this method let us imagine the following scenario: T1 chooses p1 as a candidate, while T1 is reading it, T2 reads it and sees it is already red. T2 assumes another thread is handling p1 and moves on to a bigger candidate p2 and remains ahead of T1 for the duration of the current candidate iteration.  
In this case we see that T2 has a bigger candidate and is ahead of T1 in the list.
2. In the method explained above, let us take the following scenario: T1 chooses 2 as a candidate and is reading it. T2 sees 2 is taken, moves on to 3 and is reading it. Then T3 comes, while both T1 and T2 are reading their respective candidates, and then moves on to take 4 as a candidate. Since 4 is not prime, we see a scenario where a non-prime is chosen as a candidate (although it will not remain in the list after the algorithm finishes running).

Q3:

1. First let us look at what it means to be able to change from a reader to a writer atomically. In order to be a reader and change to a writer without waiting after unlocking read and before locking write we must be the only readers and have no other writers (otherwise we risk data corruption). Thus what the heuristic of the upgrade from read to write function is: Wait till I'm the last reader, and then change to writer atomically.  
With this heuristic the problem is more easily identified. If 2 threads are allowed to upgrade from read to write, a scenario where both are readers and are waiting to be able to upgrade can occur. In this scenario both threads are waiting till they are the only readers remaining. They will wait indefinitely and we are going to have a deadlock.
2. Code given Below

```

#include <pthread.h>

int number_of_readers;
pthread_cond_t readers_condition;
int number_of_writers;
pthread_cond_t writers_condition;
mutex_t global_lock;
// bool says if a thread was selected to become the only one that can upgrade.
// upgradable stores the thread id of the thread that is upgradable.
// we use bool instead of looking for a NULL like initial value for upgradeable.
pthread_t upgradeAble;
bool upgraded;
// this condition checks when we can upgrade
pthread_cond_t R_to_W_condition;

void readers_writers_init() {
    number_of_readers = 0;
    pthread_cond_init(&readers_condition, NULL);
    number_of_writers = 0;
    pthread_cond_init(&writers_condition, NULL);
    pthread_mutex_init(&global_lock, NULL);
    upgraded = false;
}

bool upgrade_to_write_lock()
{
    pthread_mutex_lock(&global_lock);
    //if no upgrAble was chosen, choose it
    if (!upgraded){
        upgradeAble = pthread_self();
        upgraded = true;
    }
    //if the current thread is not upgradable, return 0;
    if (!pthread_equal(upgradeAble, pthread_self())){
        pthread_mutex_unlock(&global_lock);
        return false;
    }

    if (number_of_writers > 0){
        // cant be a reader since there is a writer active so either we are non
        readers or are the writer itself
        return false;
    }

    while (number_of_readers != 1 && number_of_writers != 0){
        pthread_cond_wait(&R_to_W_condition, &global_lock);
    }

    number_of_writers++;
    number_of_readers--;
    pthread_mutex_unlock(&global_lock);
    return true;
}

```

```

void read_lock() {
    pthread_mutex_lock(&global_lock);
    while (number_of_writers > 0)
        pthread_cond_wait(&readers_condition, &global_lock);
    number_of_readers++;
    pthread_mutex_unlock(&global_lock);
}

void read_unlock() {
    pthread_mutex_lock(&global_lock);
    number_of_readers--;
    if (number_of_readers == 0)
        pthread_cond_signal(&writers_condition);
    //might need to signal to upgrade if only 1 reader
    else if (number_of_readers == 1)
        pthread_cond_signal(&R_to_W_condition);

    pthread_mutex_unlock(&global_lock);
}

void write_lock() {
    pthread_mutex_lock(&global_lock);
    while ((number_of_writers > 0) || (number_of_readers > 0))
        pthread_cond_wait(&writers_condition, &global_lock);
    number_of_writers++;
    pthread_mutex_unlock(&global_lock);
}

void write_unlock() {
    pthread_mutex_lock(&global_lock);
    number_of_writers--;
    if (number_of_writers == 0) {
        if (number_of_readers == 1)
            pthread_cond_signal(&R_to_W_condition);
        pthread_cond_broadcast(&readers_condition);
        pthread_cond_signal(&writers_condition);
    }
    pthread_mutex_unlock(&global_lock);
}

```